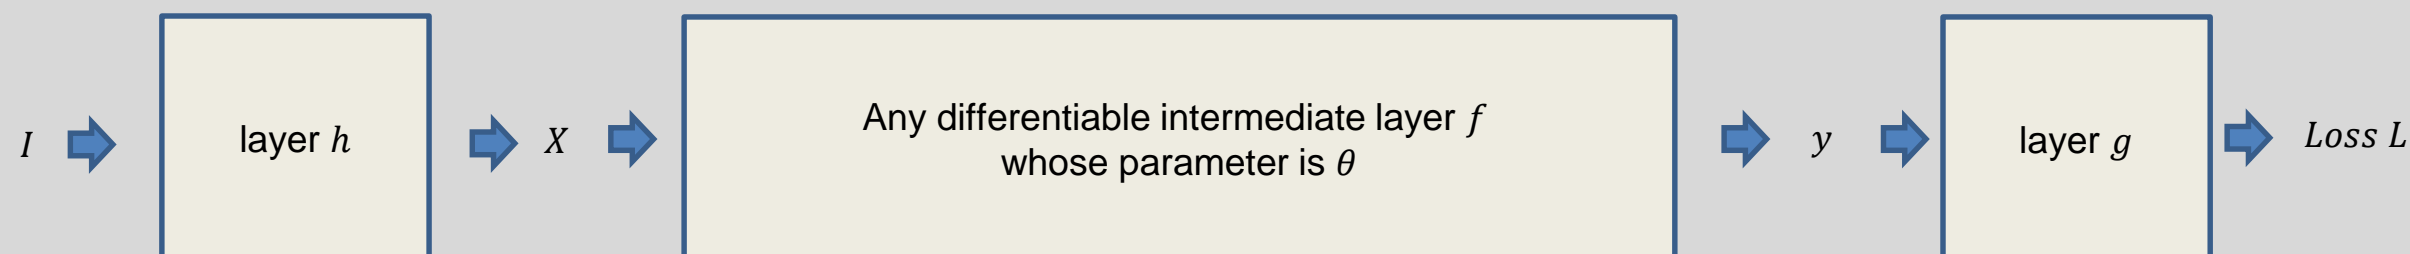


# Computer Vision

## Lecture 04: Convolutional Neural Network-1

# Differentiable layers



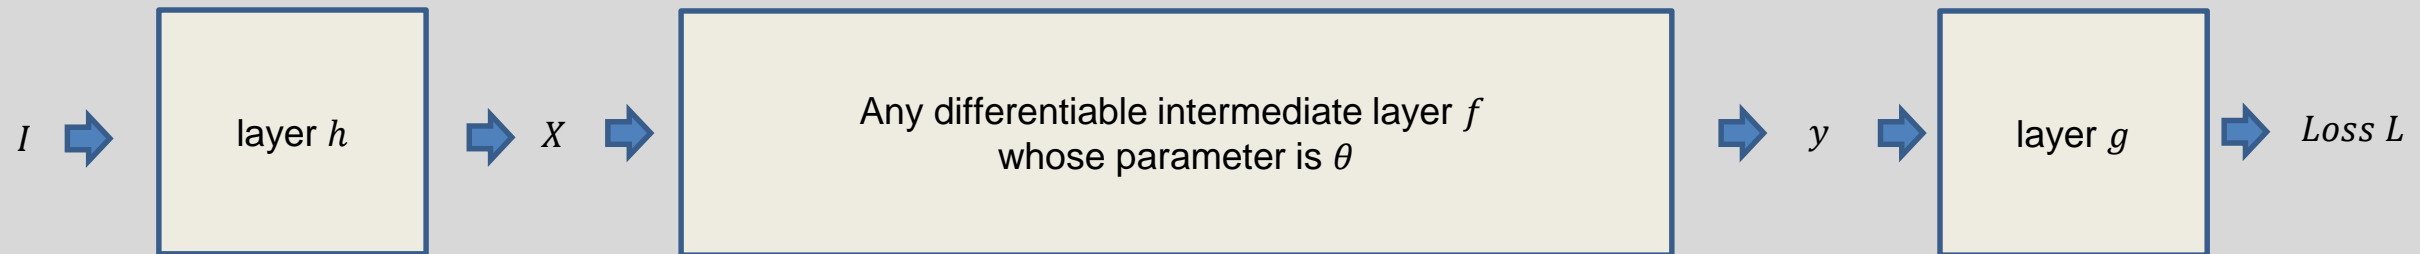
We need to implement three things for an intermediate layer  $f$  :

forward rule:  $y = f(X; \theta)$  for  $g(f(X; \theta)) = L$

backward rule:  $\frac{dy}{dX}$  for  $\frac{dL}{dX} = \frac{dy}{dX} \times \frac{dL}{dy}$

parameter update rule:  $\frac{dy}{d\theta}$  for  $\theta^{new} = \theta - \varepsilon \frac{dy}{d\theta} \times \frac{dL}{dy}$

# Differentiable layers



We need to implement three things for an intermediate layer  $f$  :

forward rule:  $y = f(X; \theta)$  for  $g(f(X; \theta)) = L$

backward rule:  $\frac{dy}{dX}$  for  $\frac{dL}{dX} = \frac{dy}{dX} \times \frac{dL}{dy}$

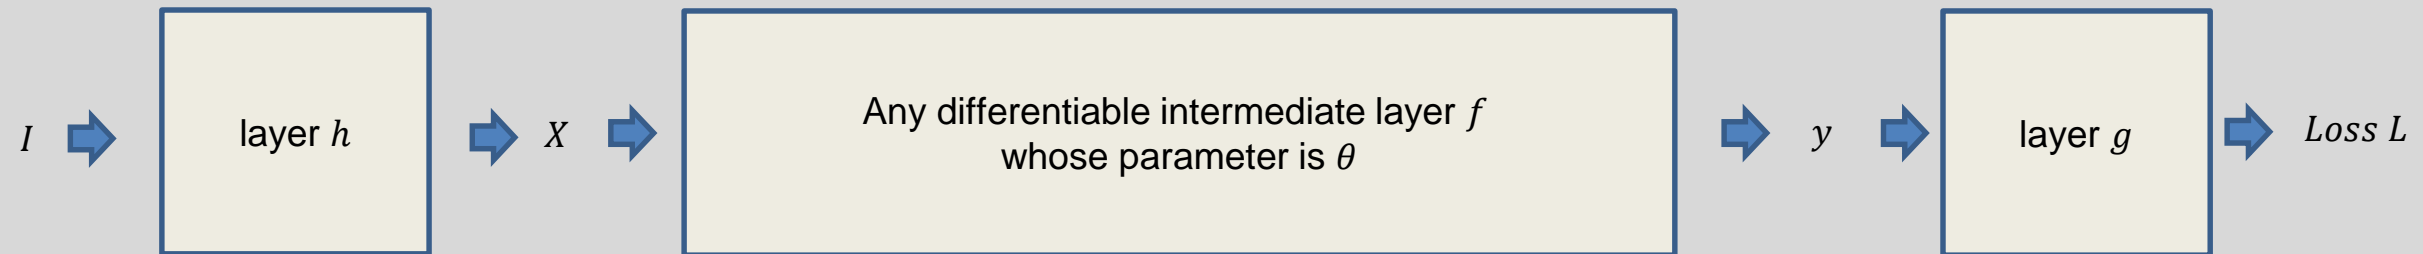
parameter update rule:  $\frac{dy}{d\theta}$  for  $\theta^{new} = \theta - \varepsilon \frac{dy}{d\theta} \times \frac{dL}{dy}$



PyTorch can do these automatically.

`loss.backward()`  
`optimizer.step()`

# Differentiable layers



Some layers (e.g. pooling, activation) do not have parameters  $\theta$ . It requires only two:

forward rule:  $y = f(X; \theta)$  for  $g(f(X; \theta)) = L$


backward rule:  $\frac{dy}{dX}$  for  $\frac{dL}{dX} = \frac{dy}{dX} \times \frac{dL}{dy}$

In convolutional layers, parameters are convolutional filter kernel's weights.

# CNNs



RGB image




cat

Semantic labels

# CNNs



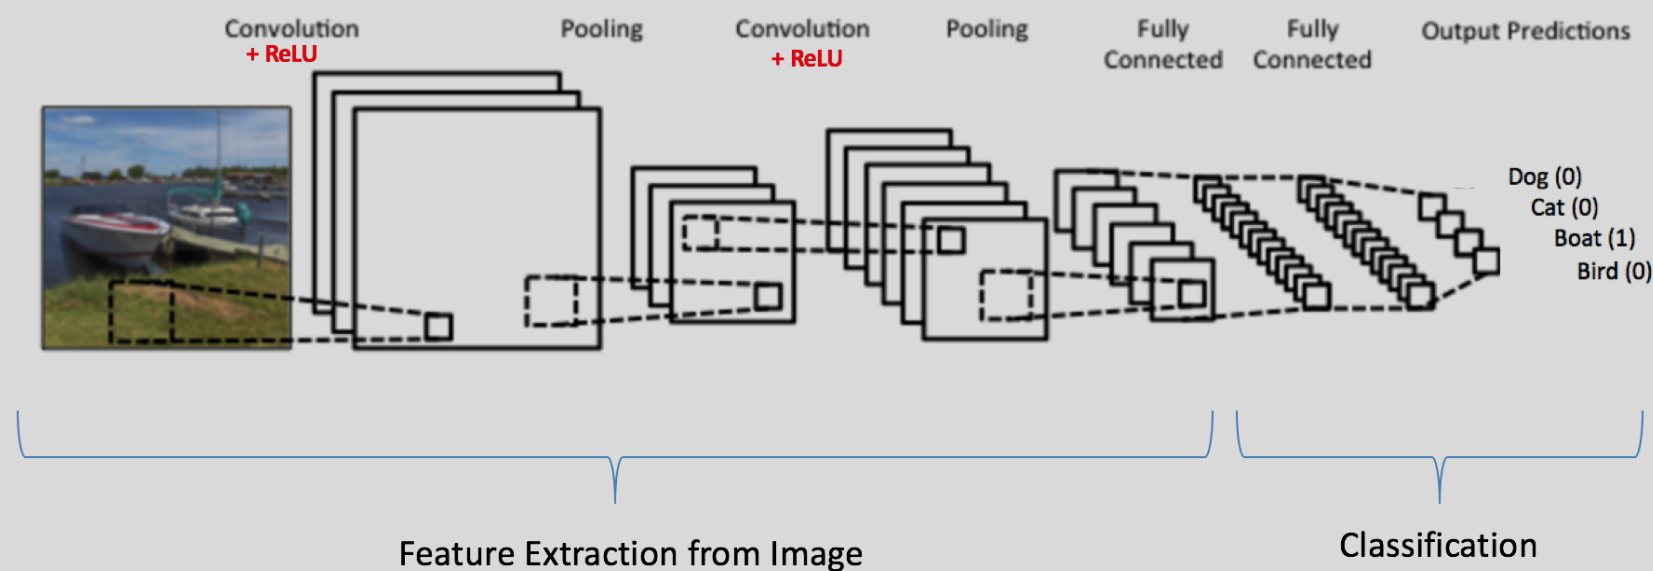
RGB image



cat

Semantic labels

# Convolutional neural network



# 2D Convolution

1	1	0	1	2
2	0	0	1	0
0	0	2	2	1
0	0	0	2	0
0	0	2	1	1

Image (5x5)

0	0	-1
1	-1	0
0	1	-1

Filter kernel (3x3)



# 2D Convolution

1	1	0	1	2
2	0	0	1	0
0	0	2	2	1
0	0	0	2	0
0	0	2	1	1

Image (5x5)

0	0	-1
1	-1	0
0	1	-1

Filter kernel (3x3)

$$\begin{aligned} &1*0+1*0+0*-1 \\ &+ 2*1+0*-1+0*0 \\ &+ 0*0+0*1+2*-1 \end{aligned}$$

0
---

# 2D Convolution

1	1	0	1	2
2	0	0	1	0
0	0	2	2	1
0	0	0	2	0
0	0	2	1	1

Image (5x5)

0	0	-1
1	-1	0
0	1	-1

Filter kernel (3x3)

$$\begin{aligned} &1*0+0*0+1*-1 \\ &+ 0*1+0*-1+1*0 \\ &+ 0*0+2*1+2*-1 \end{aligned}$$

0	-1
---	----

# 2D Convolution

1	1	0	1	2
2	0	0	1	0
0	0	2	2	1
0	0	0	2	0
0	0	2	1	1

Image (5x5)

0	0	-1
1	-1	0
0	1	-1

Filter kernel (3x3)

$$\begin{aligned} &0*0+1*0+2*-1 \\ &+ 0*1+1*-1+0*0 \\ &+ 2*0+2*1+1*-1 \end{aligned}$$

0	-1	-2
---	----	----

# 2D Convolution

1	1	0	1	2
2	0	0	1	0
0	0	2	2	1
0	0	0	2	0
0	0	2	1	1

Image (5x5)

0	0	-1
1	-1	0
0	1	-1

Filter kernel (3x3)

$$\begin{aligned} &2*0+0*0+0*-1 \\ &+ 0*1+0*-1+2*0 \\ &+ 0*0+0*1+0*-1 \end{aligned}$$

0	-1	-2
0		

# 2D Convolution

1	1	0	1	2
2	0	0	1	0
0	0	2	2	1
0	0	0	2	0
0	0	2	1	1

Image (5x5)

0	0	-1
1	-1	0
0	1	-1


Filter kernel (3x3)

$$\begin{aligned} &2*0+2*0+1*-1 \\ &+ 0*1+2*-1+2*0 \\ &+ 2*0+1*1+1*-1 \end{aligned}$$

0	-1	-2
0	-5	2
-4	-1	-3

Output feature (3x3)

# Ex. Image blurring



# 2D Convolution w/ zero padding






Image (5x5)



Filter kernel (3x3)

$$\begin{aligned}
 &0*0+0*0+0*-1 \\
 &+ 0*1+1*-1+1*0 \\
 &+ 0*0+2*1+0*-1
 \end{aligned}$$

1



# 2D Convolution w/ zero padding






Image (5x5)



Filter kernel (3x3)

$$\begin{aligned} &0*0+0*0+0*-1 \\ &+ 1*1+1*-1+0*0 \\ &+ 2*0+0*1+0*-1 \end{aligned}$$





# 2D Convolution w/ zero padding






Image (5x5)



Filter kernel (3x3)

$$\begin{aligned} &2*0+0*0+0*-1 \\ &+ 1*1+1*-1+0*0 \\ &+ 0*0+0*1+0*-1 \end{aligned}$$



Output feature (5x5)

# 2D Convolution w/ zero padding






Image (5x5)

→  $(K-1) / 2$  padding is required to obtain the original size.

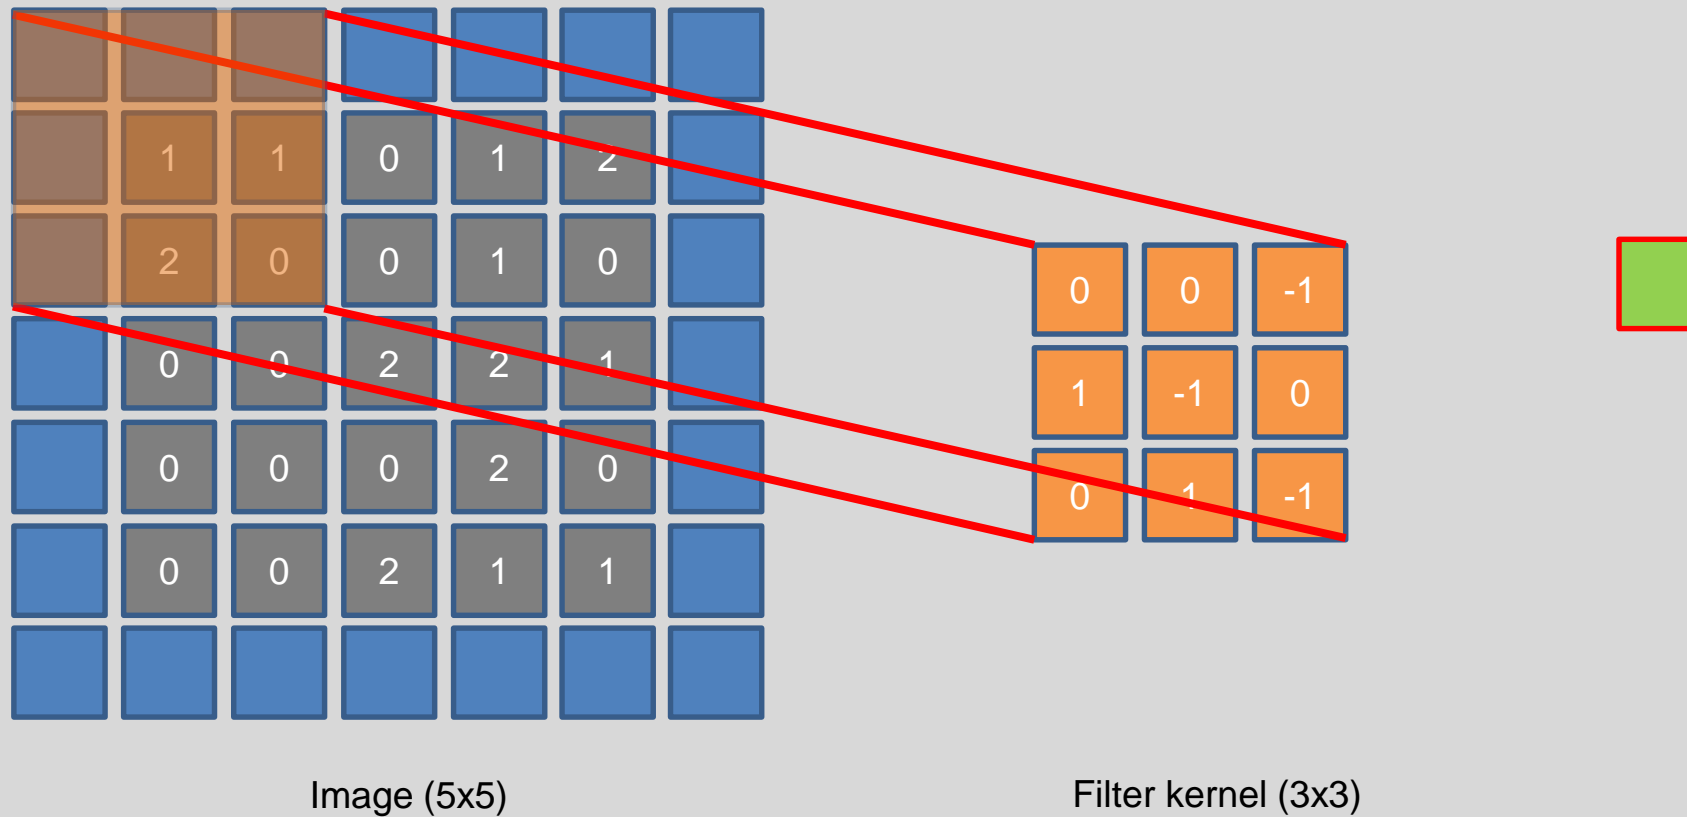


Filter kernel (5x5)



Output feature (5x5)

# 2D Convolution w/ stride



# 2D Convolution w/ stride




Image (5x5)

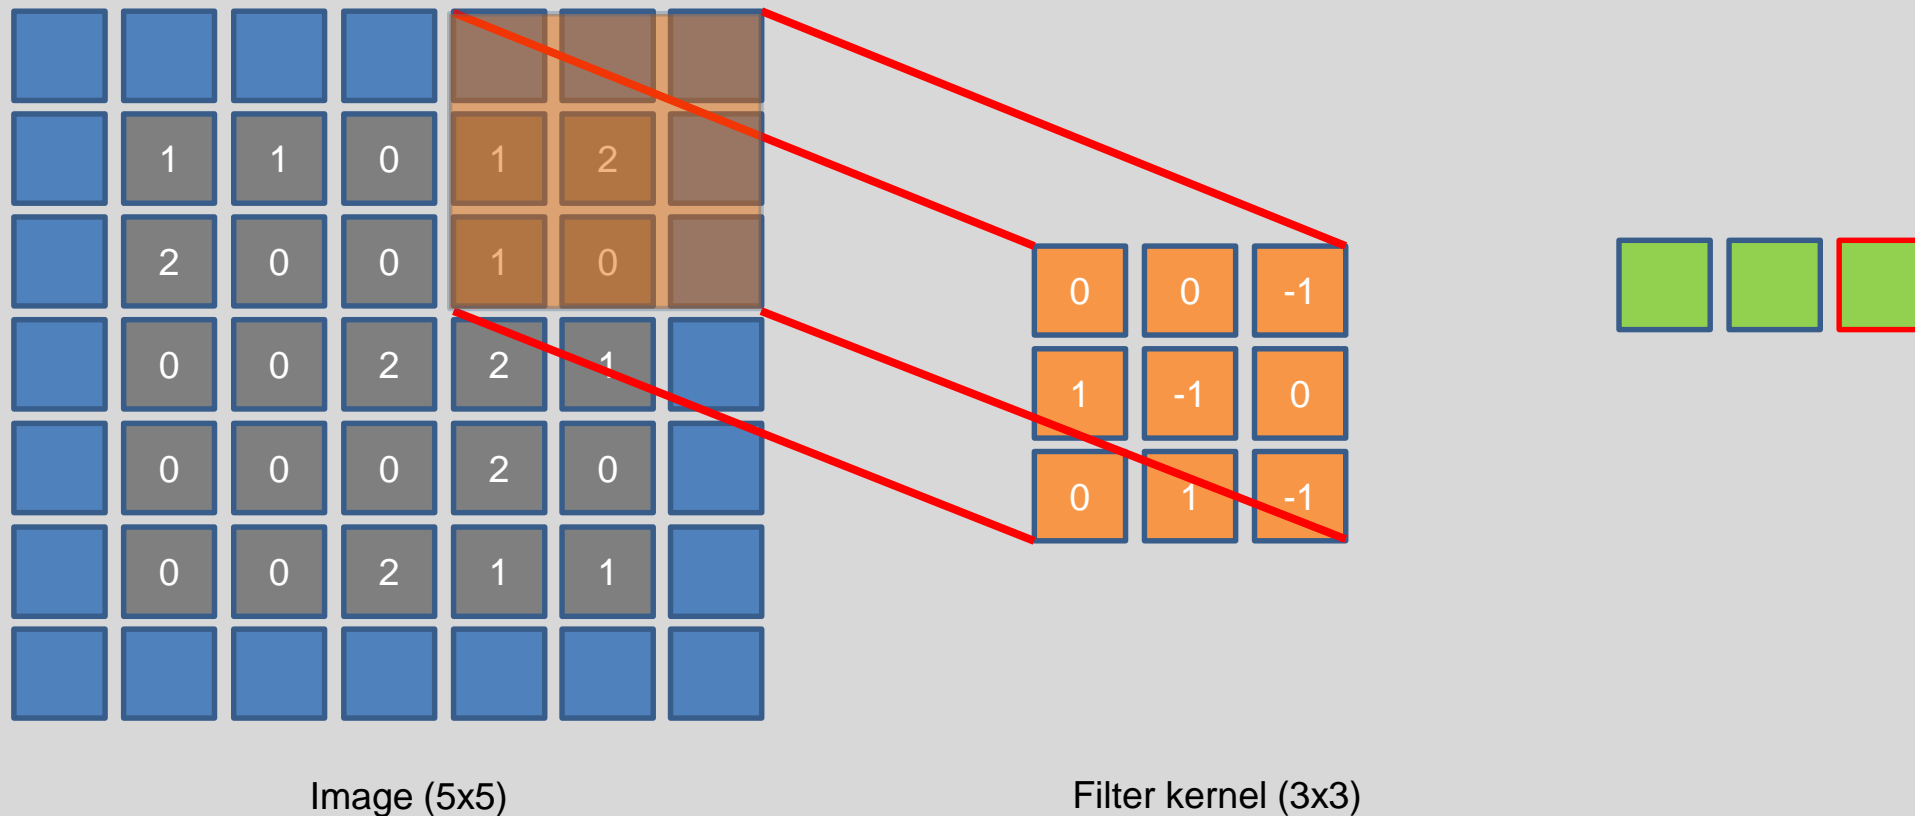


Filter kernel (3x3)

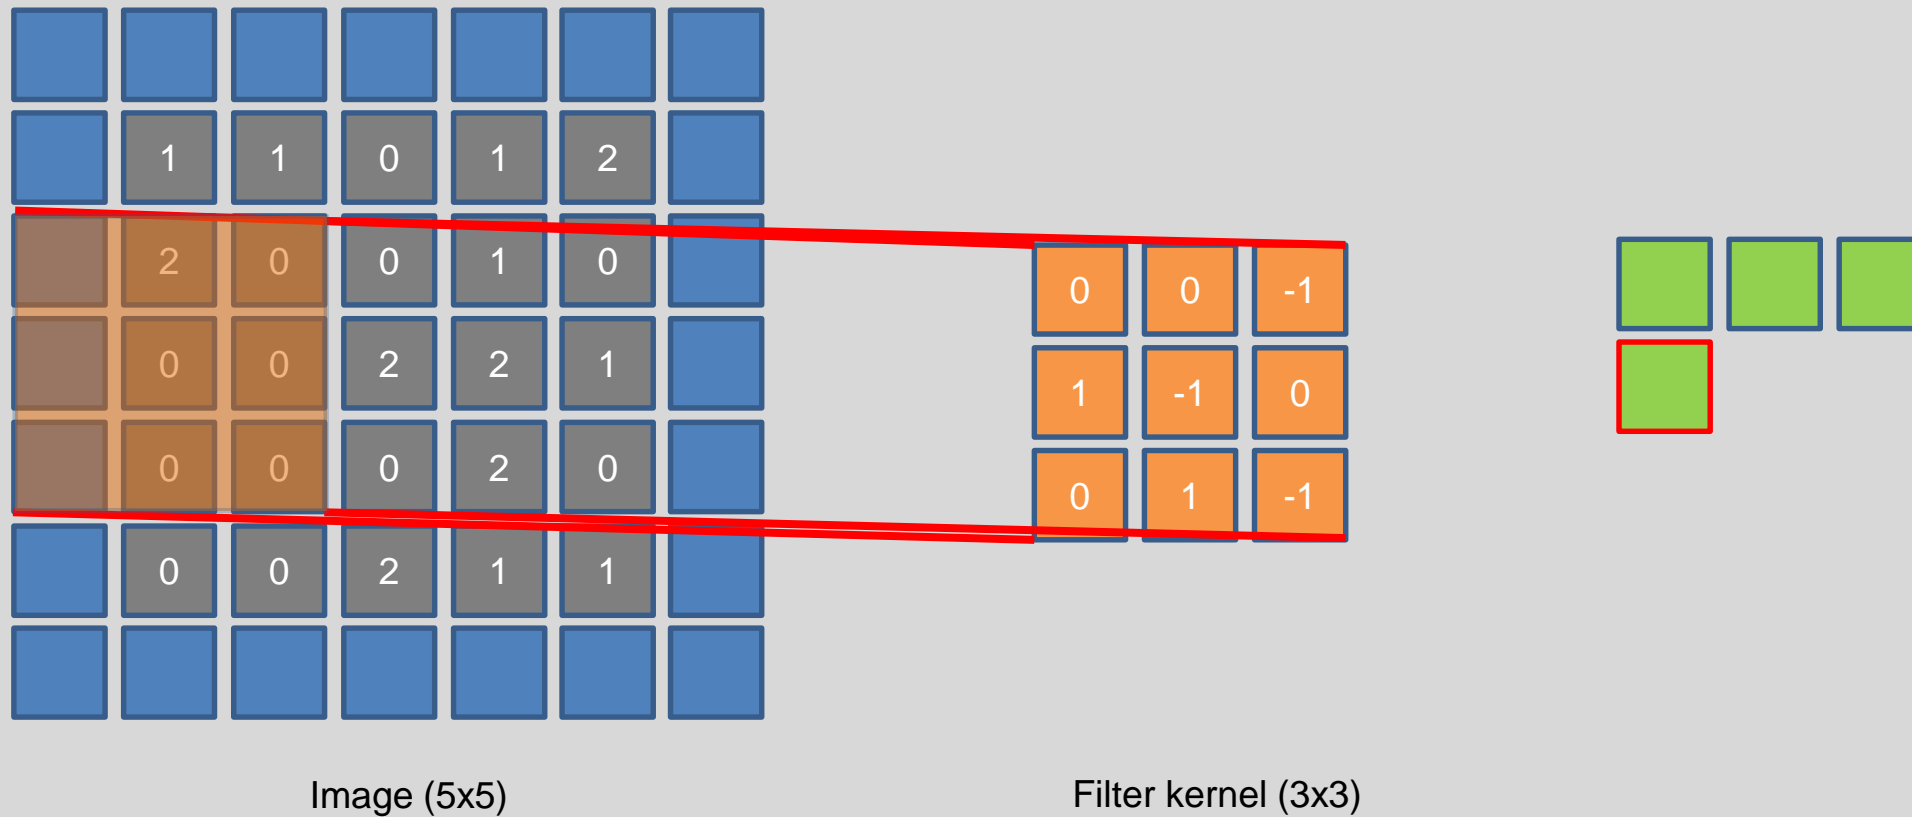
Stride=2




# 2D Convolution w/ stride




# 2D Convolution w/ stride



# 2D Convolution w/ stride



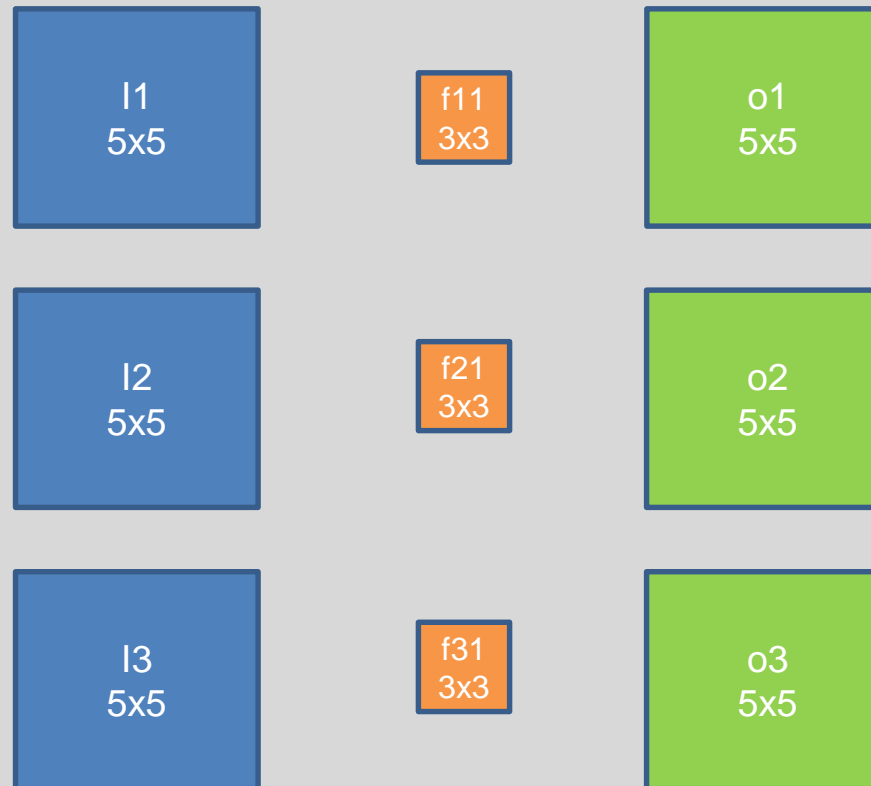
# 2D Convolution w/ stride




Stride=2



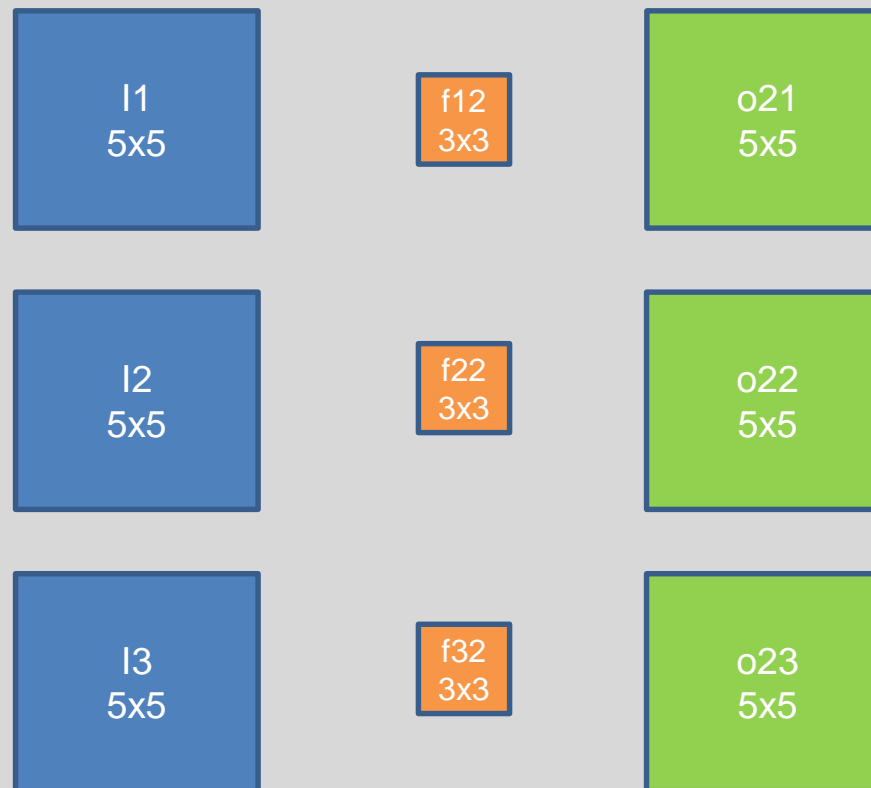
# Convolutional layer




# Convolutional layer



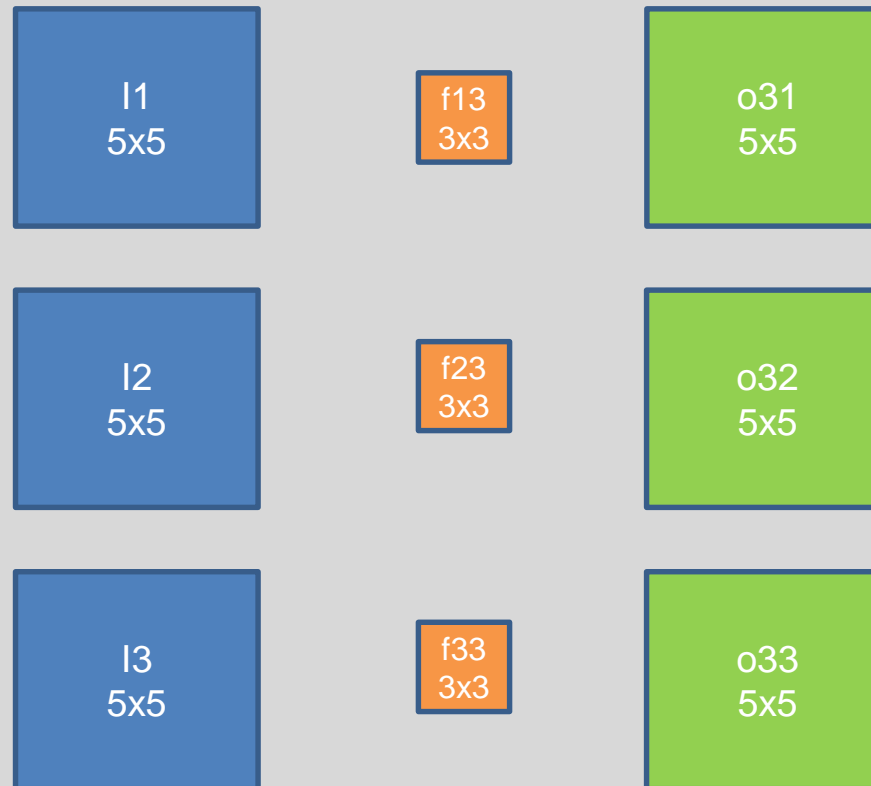
# Convolutional layer




# Convolutional layer



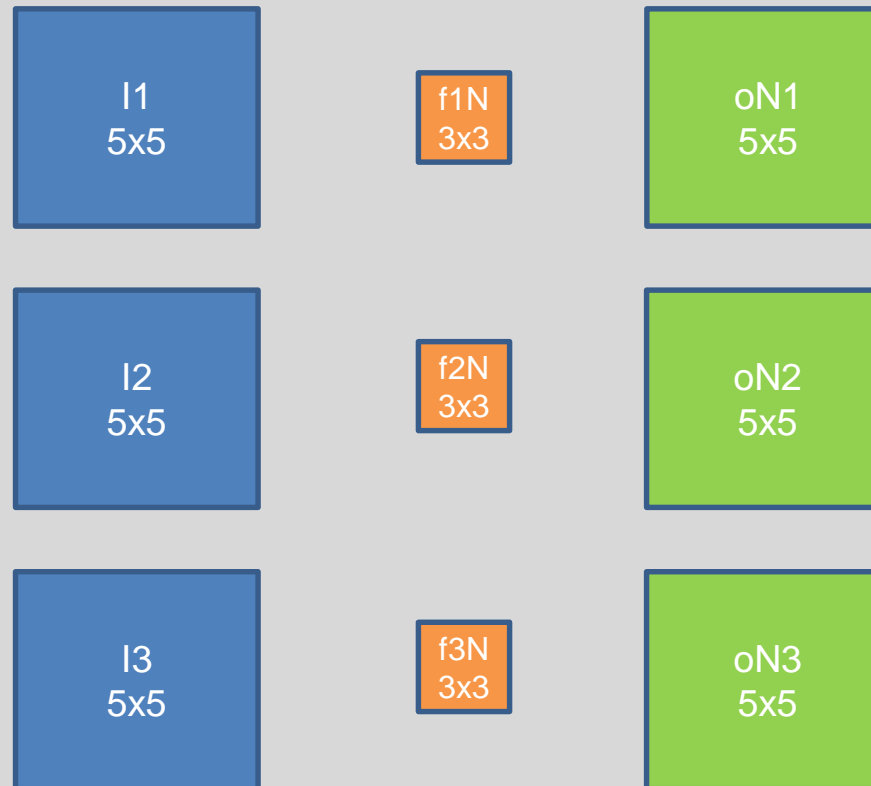
# Convolutional layer



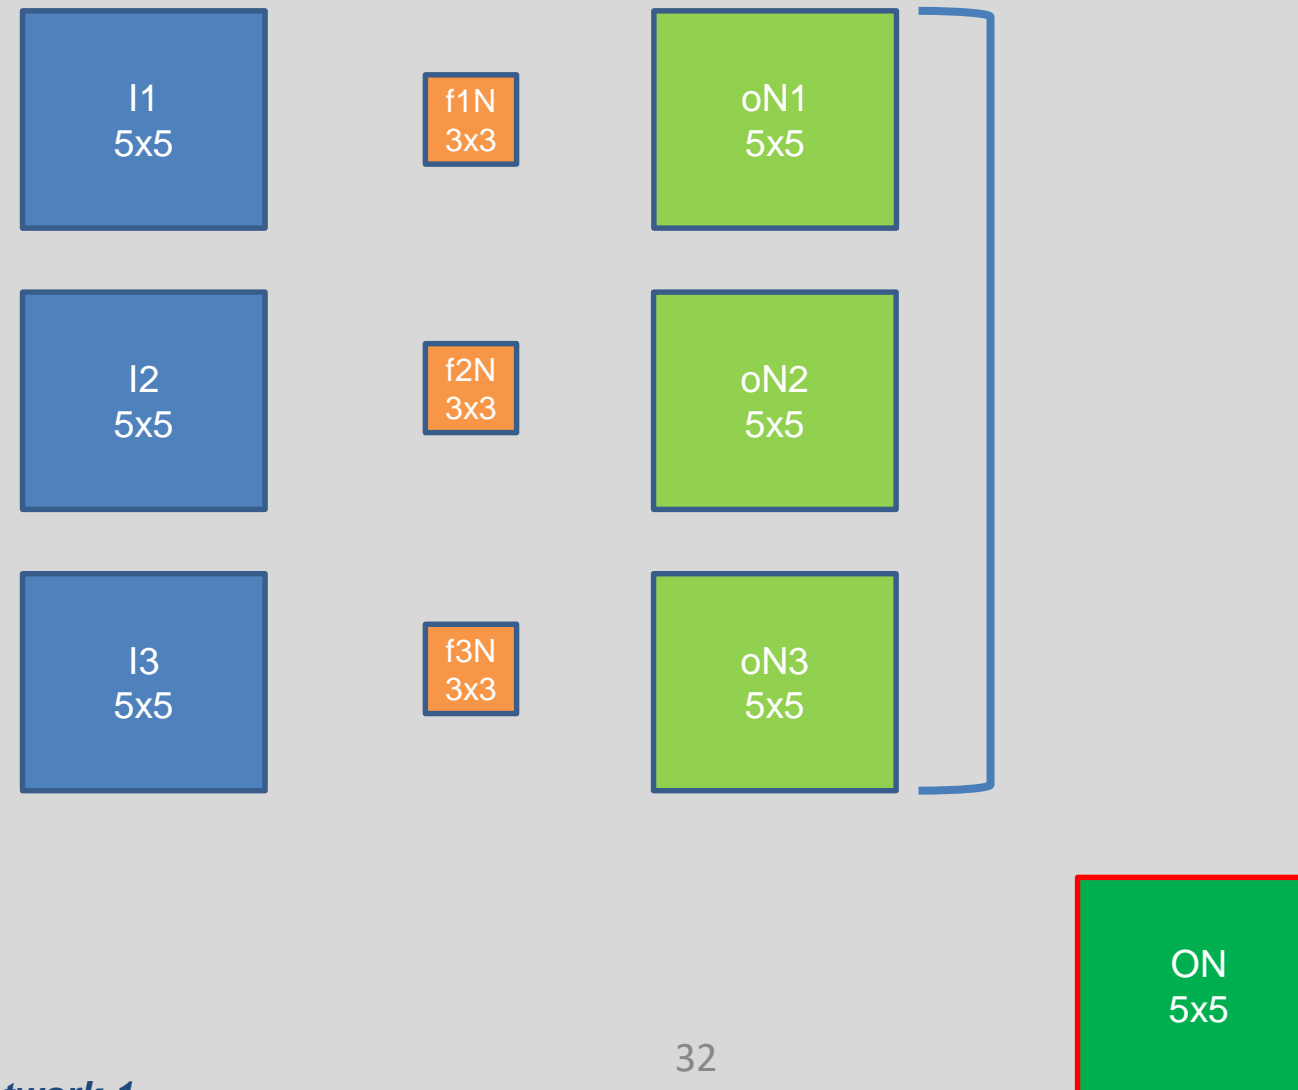
# Convolutional layer



# Convolutional layer

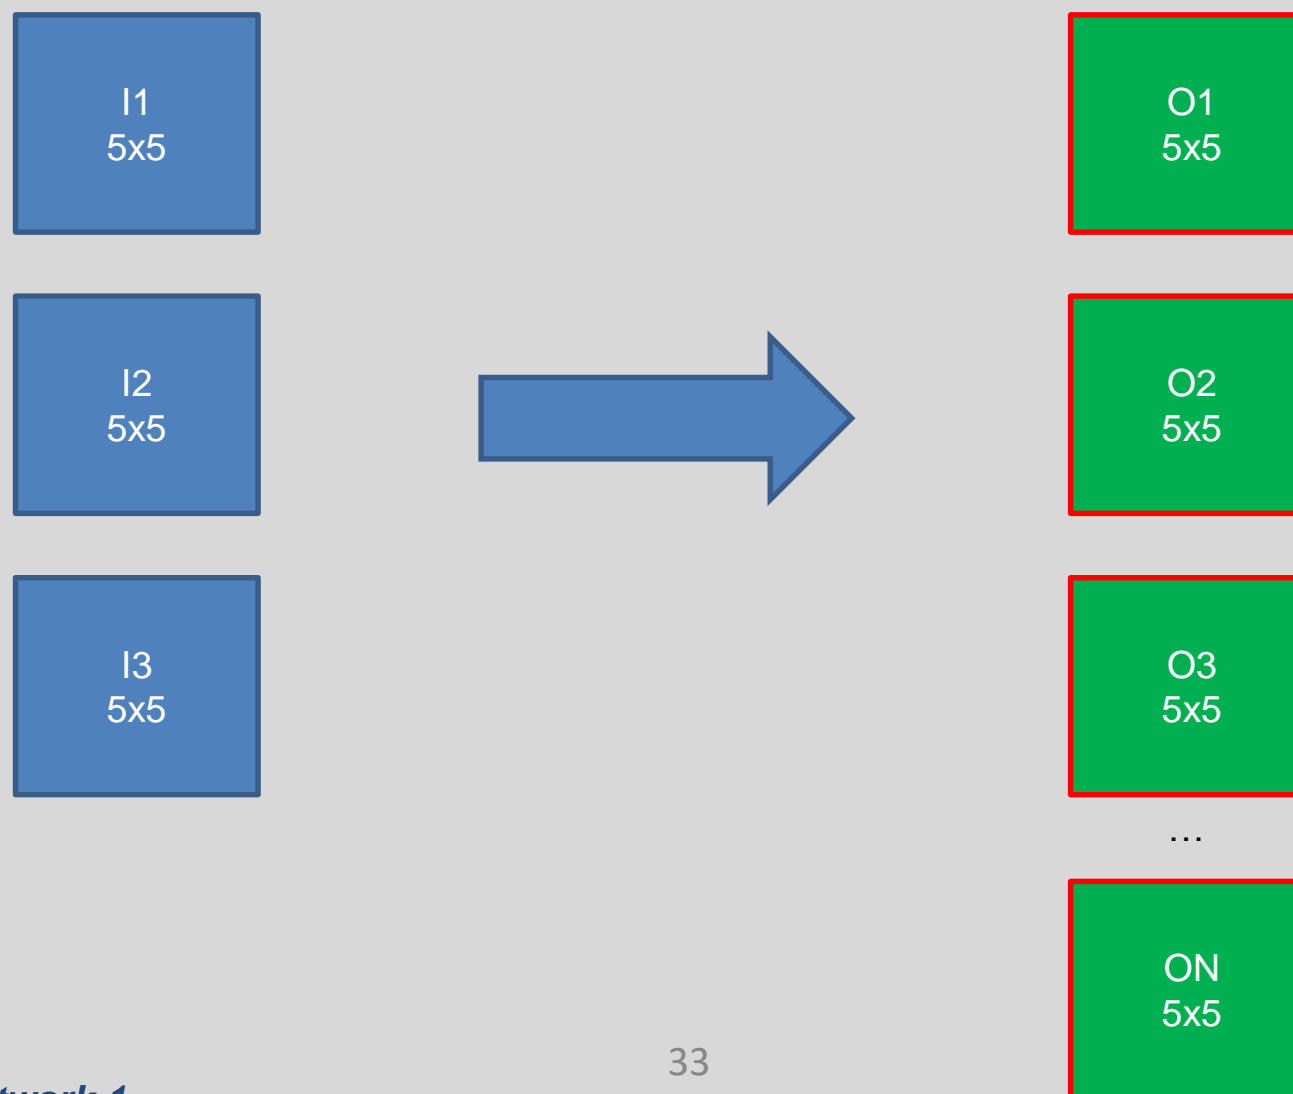


# Convolutional layer






# Convolutional layer



# Convolutional neural network



# Convolutional layer in PyTorch






Image (5x5)



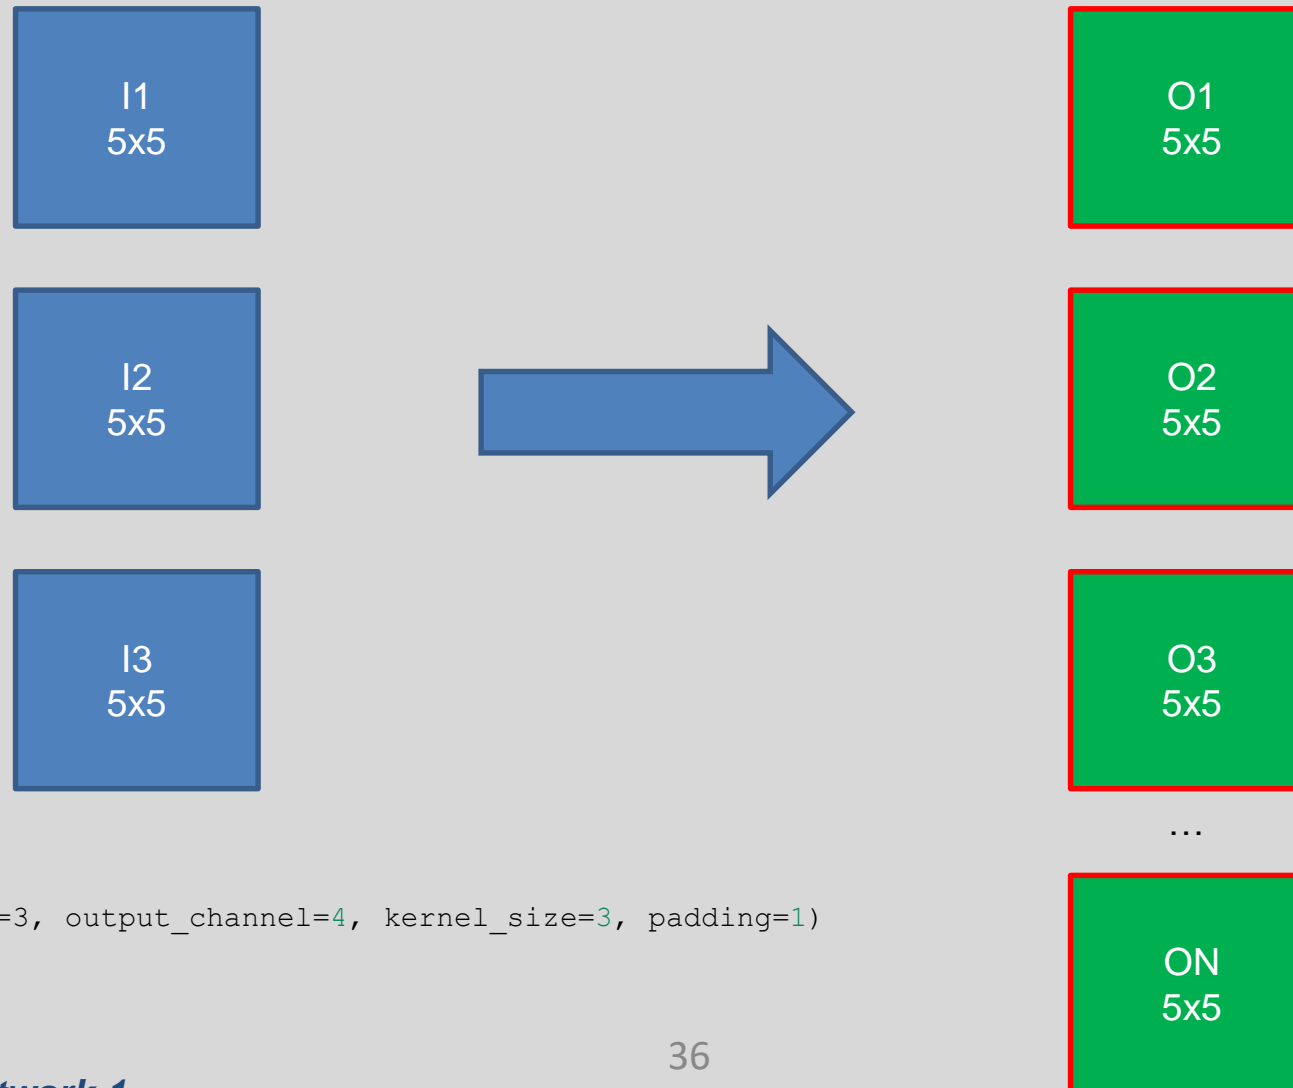
Filter kernel (3x3)



Output feature (5x5)


```
torch.nn.Conv2d(input_channel=1, output_channel=1, kernel_size=3, padding=1)
```

# Convolutional layer

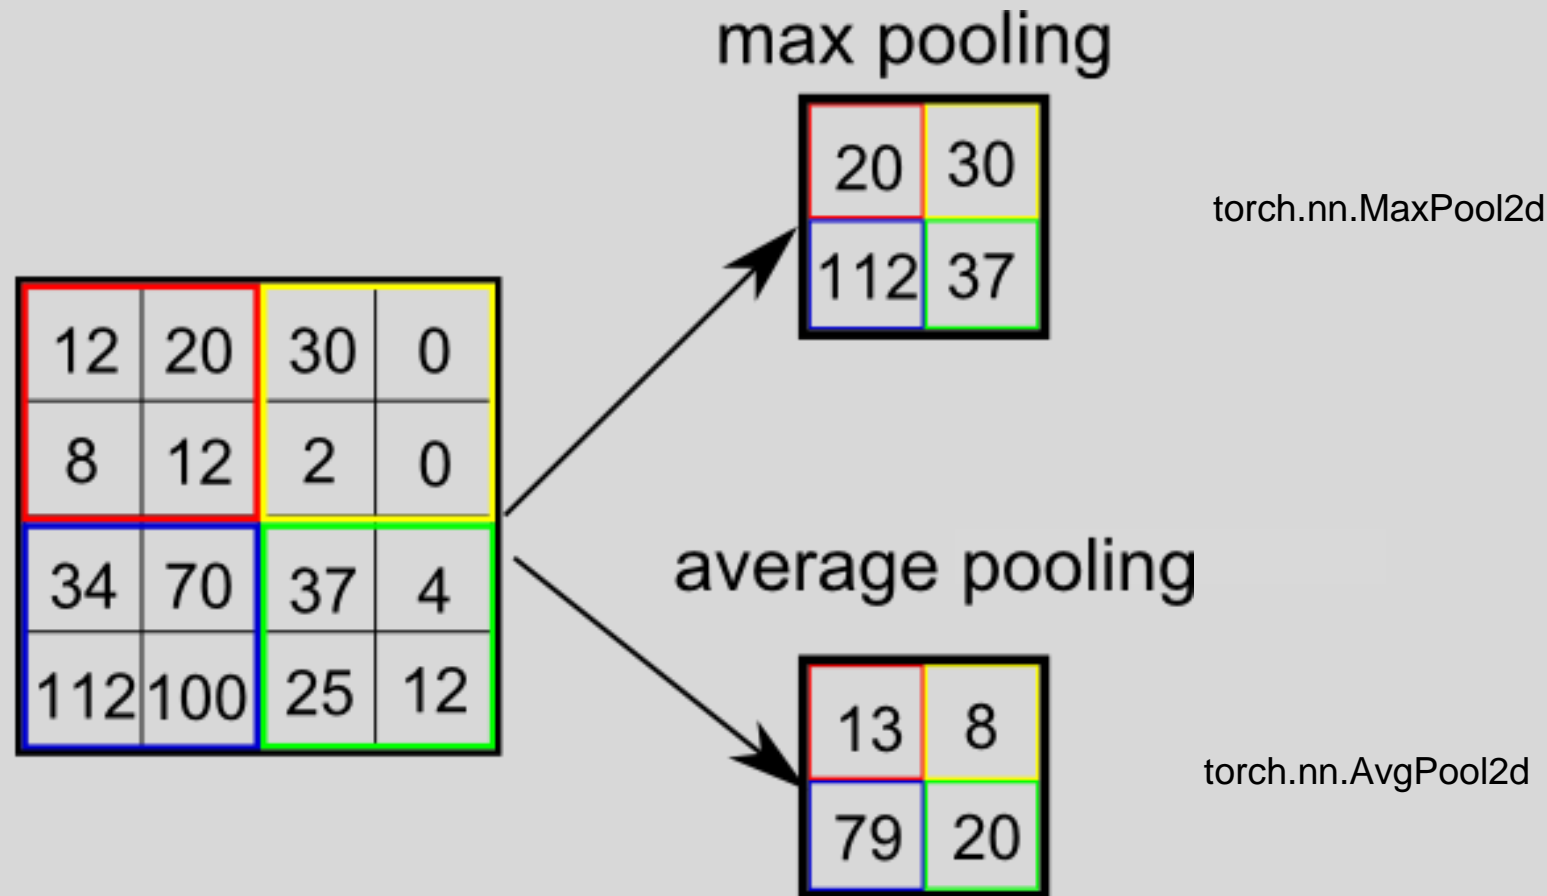


```
torch.nn.Conv2d(input_channel=3, output_channel=4, kernel_size=3, padding=1)
```

# Pooling



# Pooling




# Activation layer


$$f(x) = \frac{1}{1 + e^{-x}}$$

$$f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$


$$f(x) = \max(0, x)$$



`torch.nn.Sigmoid()`

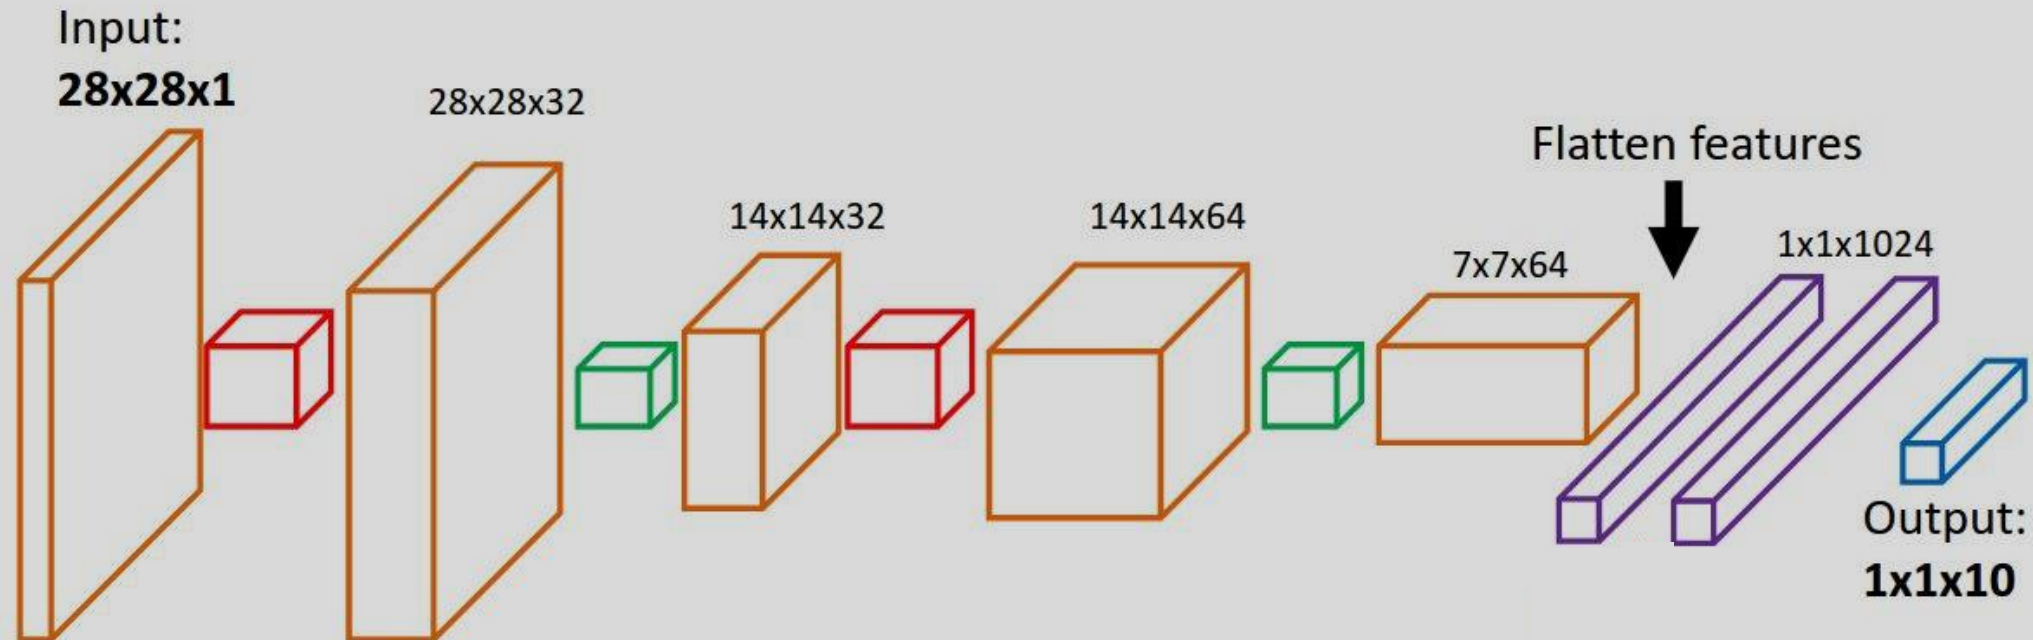


`torch.nn.Tanh()`



`torch.nn.ReLU()`

# Overall CNN architecture



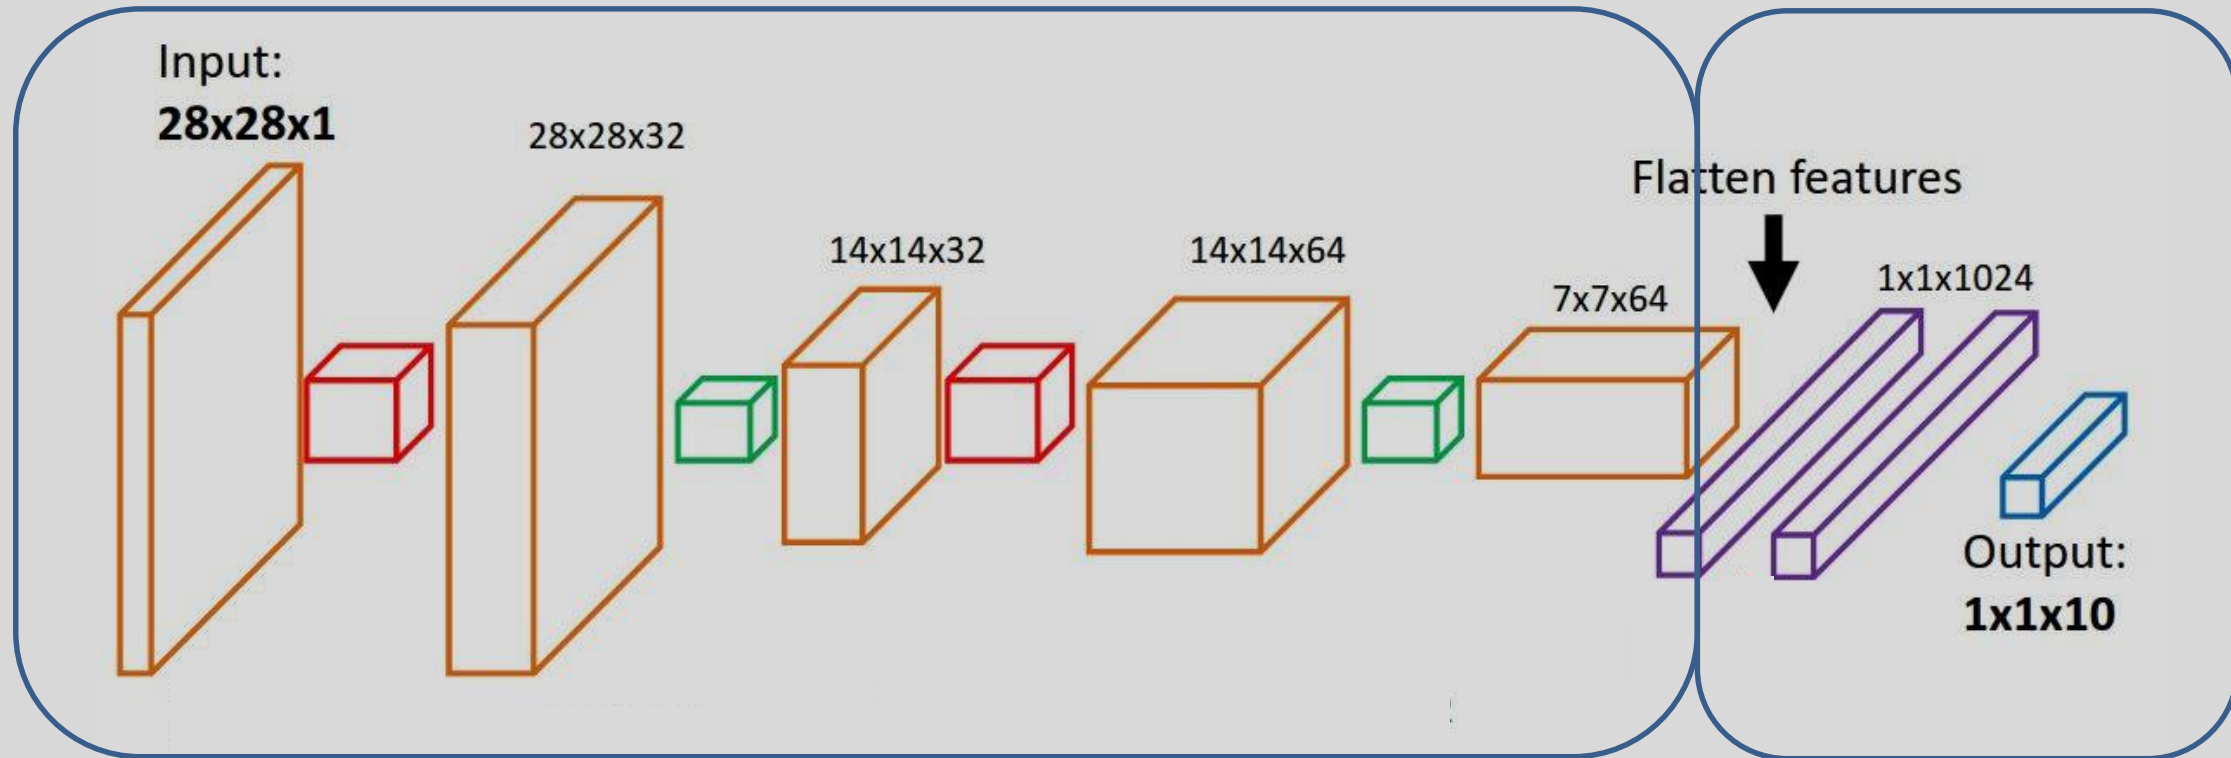
**Combination of differentiable layers → Differentiable architecture!**



# Overall CNN architecture

Convolutional layers.

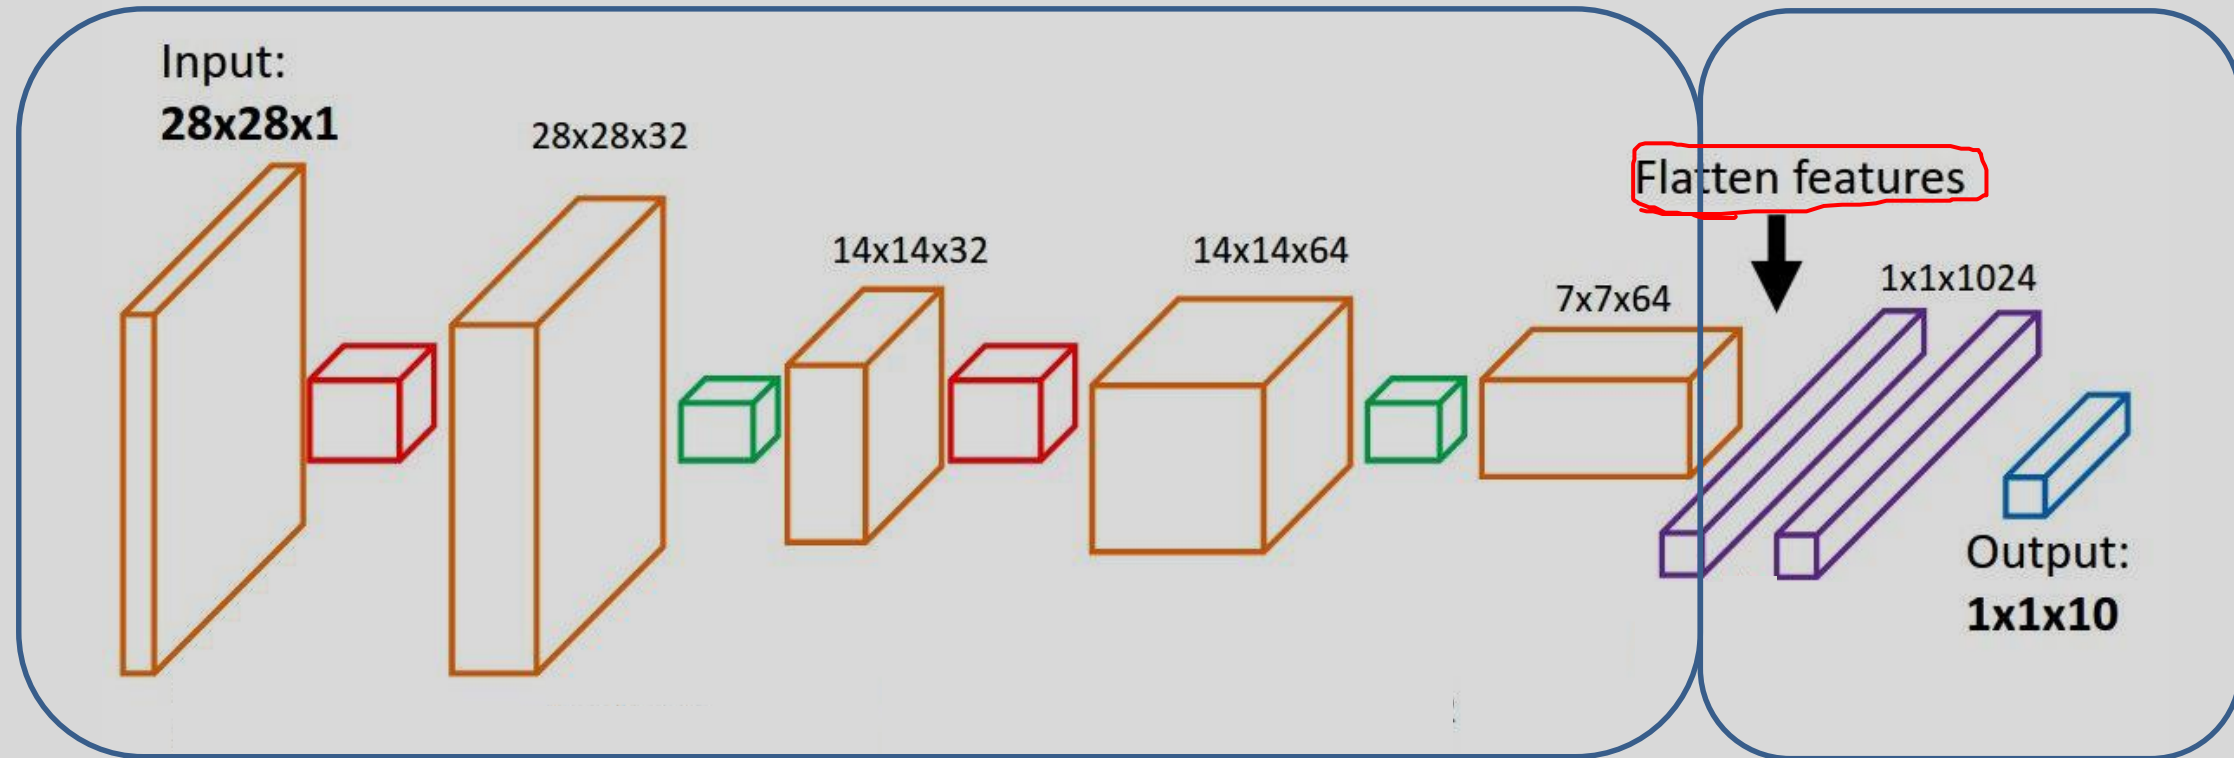
Fully-connected layers.



# Overall CNN architecture

Convolutional layers.

Fully-connected layers.




Similar to feature encoding stage of conventional machine learning (ML)  
e.g. BOW representation.

Similar to classification stage of  
conventional ML e.g. SVM.

# CNN implementation in PyTorch


```
import torch
```

```
class MyCNN(torch.nn.Module):  
    def __init__(self):  
        super().__init__()
```



Called when your network is initialized.

```
    def forward(self, x):  
        return x
```



Called when the forward pass is performed on input x.


# CNN implementation in PyTorch

```
import torch
import torch.nn


class MyCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer = nn.Sequential(
            nn.Conv2d(1, 16, 5),
            nn.ReLU(),
            nn.Conv2d(16, 32, 5),
            nn.MaxPool2d(2, 2)
            nn.Conv2d(32, 64, 5)
            nn.ReLU()
            nn.MaxPool2d(2, 2)
        )
        self.fc_layer = nn.Sequential(
            nn.Linear(64*3*3, 10)
            nn.ReLU()
            nn.Linear(100, 10)
        )

    def forward(self, x):
        out = self.layer(x)
        out = out.view(batch_size, -1)
        out = self.fc_layer(out)


        return out
```



Convolutional layers are generated.



Fully connected layers are generated.



Forward pass is defined.

Backward is automatically performed when calling `loss.backward()`

# CNN implementation in PyTorch

```
import torch
import torch.nn

class MyCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer = nn.Sequential(
            nn.Conv2d(1, 16, 5, padding=2),
            nn.ReLU(),
            nn.Conv2d(16, 32, 5),
            nn.MaxPool2d(2, 2),
            nn.Conv2d(32, 64, 5),
            nn.ReLU(),
            nn.MaxPool2d(2, 2)
        )
        self.fc_layer = nn.Sequential(
            nn.Linear(64*3*3, 10),
            nn.ReLU(),
            nn.Linear(100, 10)
        )

    def forward(self, x):
        out = self.layer(x)
        out = out.view(batch_size, -1)
        out = self.fc_layer(out)

        return out
```

```
import torch

net = MyCNN()

loss_func = torch.nn.MSELoss()
optimizer = torch.optim.SGD(net.parameters(), lr=0.01)

losses = []

for i in range(num_epoch):
    optimizer.zero_grad()

    output = net(x)

    loss = loss_func(output, y)
    loss.backward()

    optimizer.step()

    losses.append(loss.item())
```

# CNN implementation in PyTorch

```
import torch
import torch.nn

class MyCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer = nn.Sequential(
            nn.Conv2d(1, 16, 5),
            nn.ReLU(),
            nn.Conv2d(16, 32, 5),
            nn.MaxPool2d(2, 2),
            nn.Conv2d(32, 64, 5),
            nn.ReLU(),
            nn.MaxPool2d(2, 2)
        )
        self.fc_layer = nn.Sequential(
            nn.Linear(64*3*3, 10),
            nn.ReLU(),
            nn.Linear(100, 10)
        )

    def forward(self, x):
        out = self.layer(x)
        out = out.view*(batch_size, -1)
        out = self.fc_layer(out)

        return out
```

```
import torch

net = MyCNN()

loss_func = torch.nn.MSELoss()
optimizer = torch.optim.SGD(net.parameters(), lr=0.01)

losses = []

for i in range(num_epoch):
    optimizer.zero_grad()

    output = net(x)

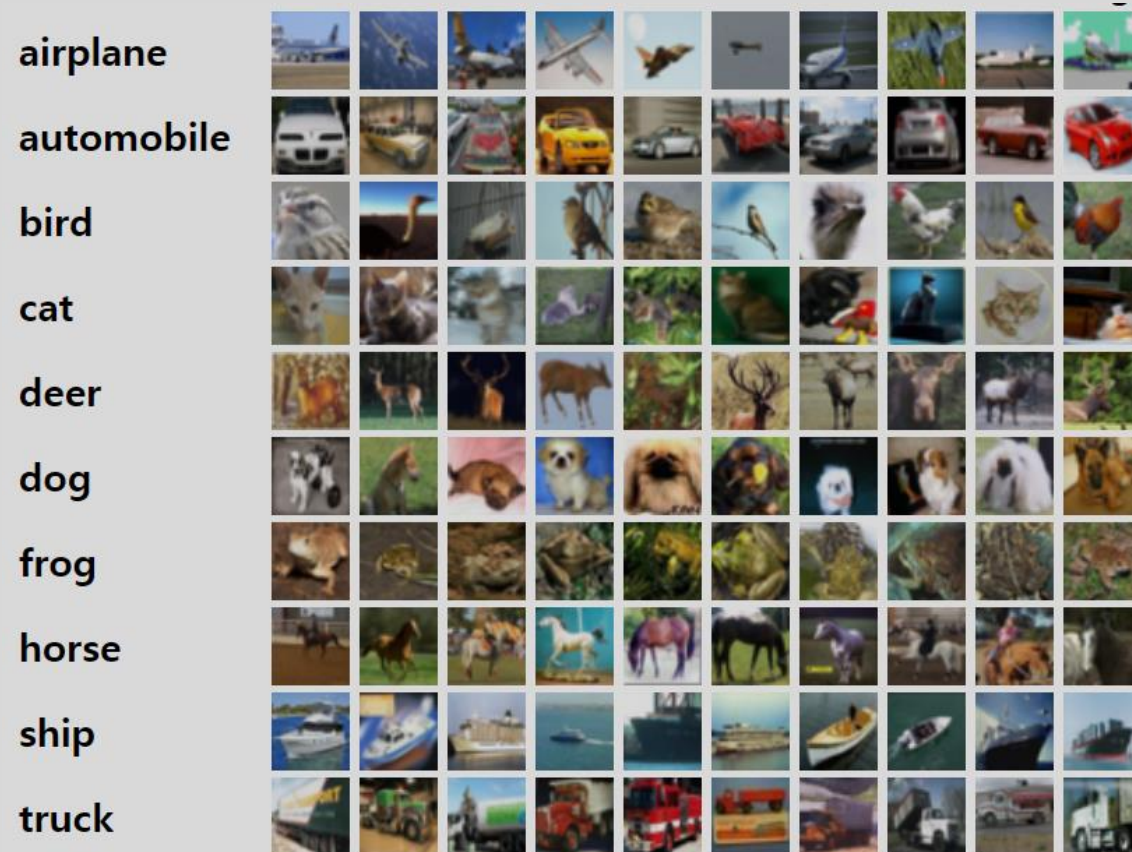
    loss = loss_func(output, y)

    loss.backward()
    optimizer.step()

    losses.append(loss.item())
```

# Image Classification task

Cifar 10 dataset (10 classes)



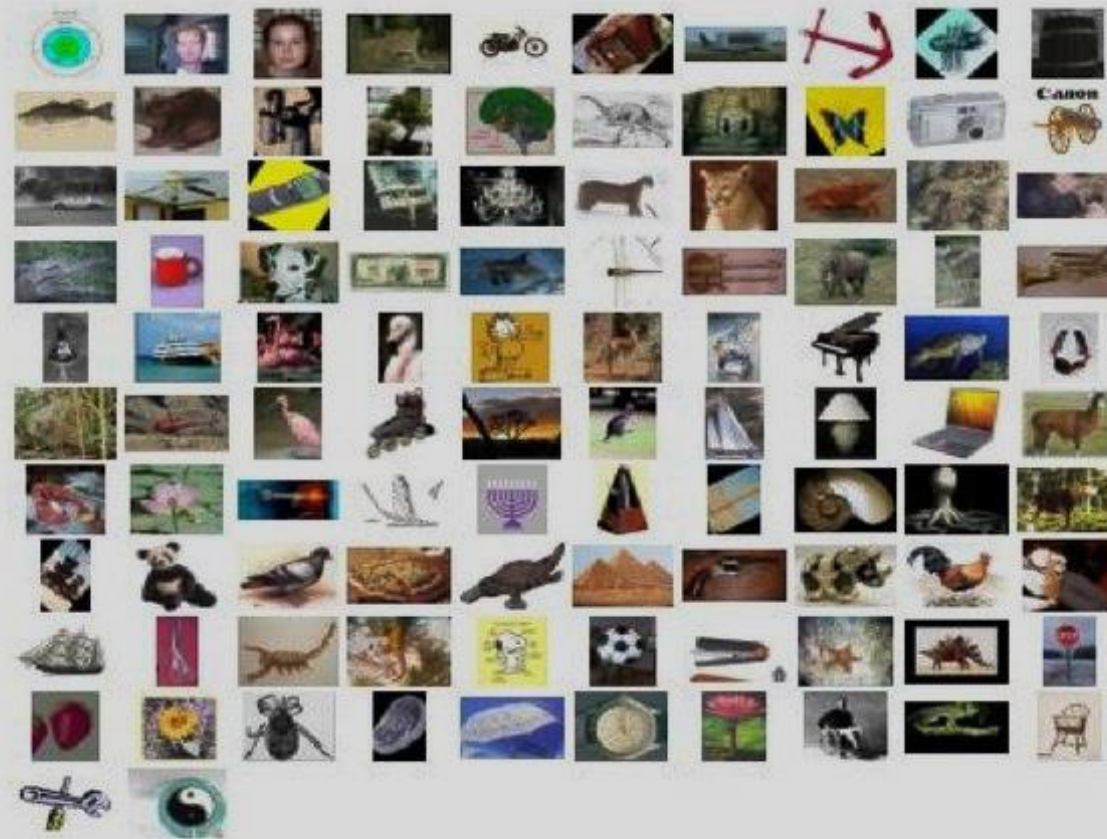
MNIST dataset (10 classes)





# Image Classification task

Caltech 101 dataset (101 classes)




ImageNet dataset (1000 classes)





# Image Classification task



# Regression Loss

```
import torch
import torch.nn as nn

loss1 = nn.MSELoss()
loss2 = nn.L1Loss()
```


$$Loss1 = ||x - y||_2^2$$

$$Loss2 = ||x - y||_1$$

# Logistic Regression

Called as the regression, but actually performs the **binary classification**!

$$\begin{aligned} z &= \frac{1}{1 + e^{-\mathbf{w}x + b}} \\ &= \sigma(-\mathbf{w}x + b) \end{aligned}$$



# Logistic Regression

$$Loss = \begin{cases} -\ln z_n, y_n = 1 \\ -\ln(1 - z_n), y_n = 0 \end{cases}$$

$$Loss = -\sum_n y_n \ln z_n + (1 - y_n) \ln(1 - z_n)$$

# Logistic Regression

```
import torch
import torch.nn as nn
from sklearn.datasets import load_iris
iris = load_iris()

X = iris.data[:100]
y = iris.target[:100]
X = torch.tensor(X, dtype=torch.float32)
y = torch.tensor(y, dtype=torch.float32)

net = nn.Linear(4, 1)
loss_fn = nn.BCELoss()
optimizer = torch.optim.SGD(net.parameters(), lr=0.1)

losses = []

for epoch in range(100):
    optimizer.zero_grad()

    h = net(X)
    prob = nn.functional.sigmoid(h)

    loss = loss_fn(prob, y)
    loss.backward()

    optimizer.step()
    losses.append(loss.item())
```

# Logistic Regression

```
import torch
import torch.nn as nn
from sklearn.datasets import load_iris
iris = load_iris()

X = iris.data[:100]
y = iris.target[:100]
X = torch.tensor(X, dtype=torch.float32)
y = torch.tensor(y, dtype=torch.float32)

net = nn.Linear(4, 1)
loss_fn = nn.BCEWithLogitsLoss()
optimizer = torch.optim.SGD(net.parameters(), lr=0.1)

losses = []

for epoch in range(100):
    optimizer.zero_grad()

    h = net(X)

    loss = loss_fn(h.view_as(y), y)
    loss.backward()

    optimizer.step()
    losses.append(loss.item())
```

# Logistic Regression

```
import torch
import torch.nn as nn
from sklearn.datasets import load_iris
iris = load_iris()

X = iris.data[[1, 51]]
y = iris.target[[1, 51]]
X = torch.tensor(X, dtype=torch.float32)
y = torch.tensor(y, dtype=torch.float32)

net = nn.Linear(4, 1)

h=net(X)
prob = nn.functional.sigmoid(h)

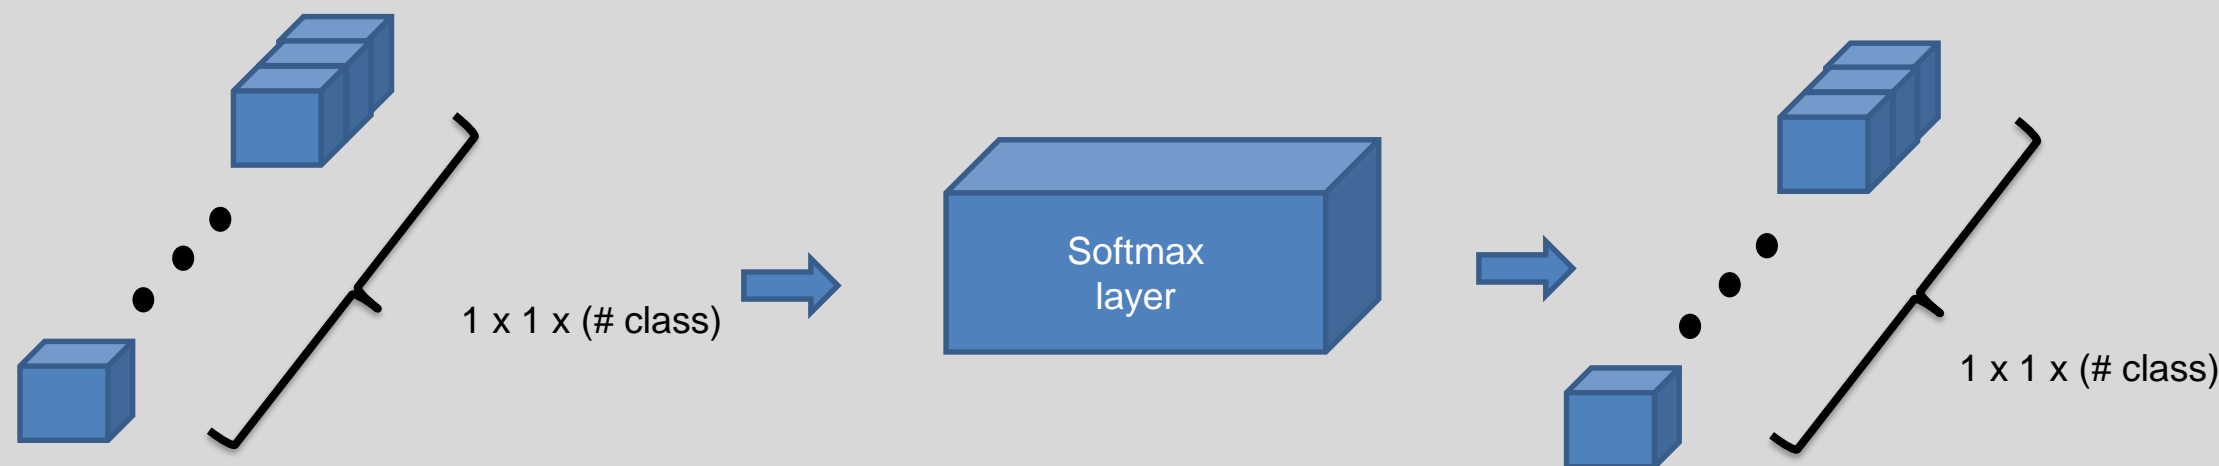
loss_fn = nn.BCELoss()
loss_fn2 = nn.BCEWithLogitsLoss()

loss1 = loss_fn(prob, y)
loss2 = loss_fn2(h.view_as(y), y)

print(loss1, loss2)
```

```
tensor(1.2881, grad_fn=<BinaryCrossEntropyBackward>) tensor(1.2881, grad_fn=<BinaryCrossEntropyWithLogitsBackward>)
```

# Softmax for multi-class classification



$$\textit{Softmax}(y_i) = \frac{\exp(y_i)}{\sum_j \exp(y_j)}$$

The vector is L1-normalized. → It could mean probability for semantic classes.



# Cross-entropy Loss

$$\begin{aligned} H(p, q) &= - \sum_x p(x) \log q(x) \\ &= \underbrace{- \sum_x p(x) \log p(x)}_{\text{Constant}} + \underbrace{\sum_x p(x) \log \frac{p(x)}{q(x)}}_{\text{KL divergence}} \end{aligned}$$

# Cross-entropy Loss

$$D_{KL}(p||q) = \sum_x p(x) \log \frac{p(x)}{q(x)}$$

The value becomes 0 when  $p(x) = q(x)$ . It is the minimum.

# Cross-entropy Loss

```
import torch
import torch.nn as nn

loss = nn.CrossEntropyLoss()


input = torch.randn(3, 5, requires_grad=True)
target = torch.empty(3, dtype=torch.long).random_(5)

output = loss(input, target)

print(input)
print(target)
print(output)
```

```
tensor([[ -1.8600, -1.1599,  0.0525, -1.9408, -1.4070],
        [ 0.6205, -1.8228,  2.1522, -2.5549, -0.1288],
        [ 0.1473, -1.5891,  1.0388, -0.6910,  0.4188]], requires_grad=True)
tensor([2, 1, 0])
tensor(2.1822, grad_fn=<NllLossBackward>)
```

# LeNet



[Yann Lecun 1998.]

# Implementing LeNet

```
import numpy as np

import torch
import torch.nn as nn

from torch.utils.data import DataLoader
from torchvision import datasets, transforms

import matplotlib.pyplot as plt

LEARNING_RATE = 0.001
BATCH_SIZE = 32
N_EPOCHS = 100

IMG_SIZE = 32
N_CLASSES = 10
```

# LeNet

```
class LeNet5(nn.Module):

    def __init__(self, n_classes):
        super(LeNet5, self).__init__()

        self.feature_extractor = nn.Sequential(
            nn.Conv2d(in_channels=1, out_channels=6, kernel_size=5, stride=1),
            nn.Tanh(),
            nn.AvgPool2d(kernel_size=2, stride=2),
            nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5, stride=1),
            nn.Tanh(),
            nn.AvgPool2d(kernel_size=2, stride=2),
            nn.Conv2d(in_channels=16, out_channels=120, kernel_size=5, stride=1),
            nn.Tanh()
        )

        self.classifier = nn.Sequential(
            nn.Linear(in_features=120, out_features=84),
            nn.Tanh(),
            nn.Linear(in_features=84, out_features=n_classes),
        )

    def forward(self, x):
        x = self.feature_extractor(x)
        x = torch.flatten(x, 1)
        x = self.classifier(x)
        return x
```

# Data Loader

```
trans = transforms.Compose([transforms.Resize((32, 32)), transforms.ToTensor()])

train_dataset = datasets.MNIST(root = 'mnist_data', train=True, transform=trans, download=True)
test_dataset = datasets.MNIST(root = 'mnist_data', train=False, transform=trans)

train_loader = DataLoader(dataset=train_dataset, batch_size=BATCH_SIZE, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=BATCH_SIZE, shuffle=False)
```

Frequently used Datasets are readily available!

<https://pytorch.org/docs/stable/torchvision/datasets.html>



# Data Loader

```
trans = transforms.Compose([transforms.Resize((32, 32)), transforms.ToTensor()])

train_dataset = datasets.MNIST(root = 'mnist_data', train=True, transform=trans, download=True)
test_dataset = datasets.MNIST(root = 'mnist_data', train=False, transform=trans)

train_loader = DataLoader(dataset=train_dataset, batch_size=BATCH_SIZE, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=BATCH_SIZE, shuffle=False)
```

You can also define your own dataset for data loader.

```
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
```

```
class CustomDataset(Dataset):
```

```
    def __init__(self):
        self.x_data = [[73, 80, 75], [93, 88, 93], [89, 91, 90], [96, 98, 100], [73, 66, 70]]
        self.y_data = [[152], [185], [180], [196], [142]]
```

```
    def __len__(self):
        return len(self.x_data)
```

```
    def __getitem__(self, idx):
        x = torch.FloatTensor(self.x_data[idx])
        y = torch.FloatTensor(self.y_data[idx])
        return x, y
```



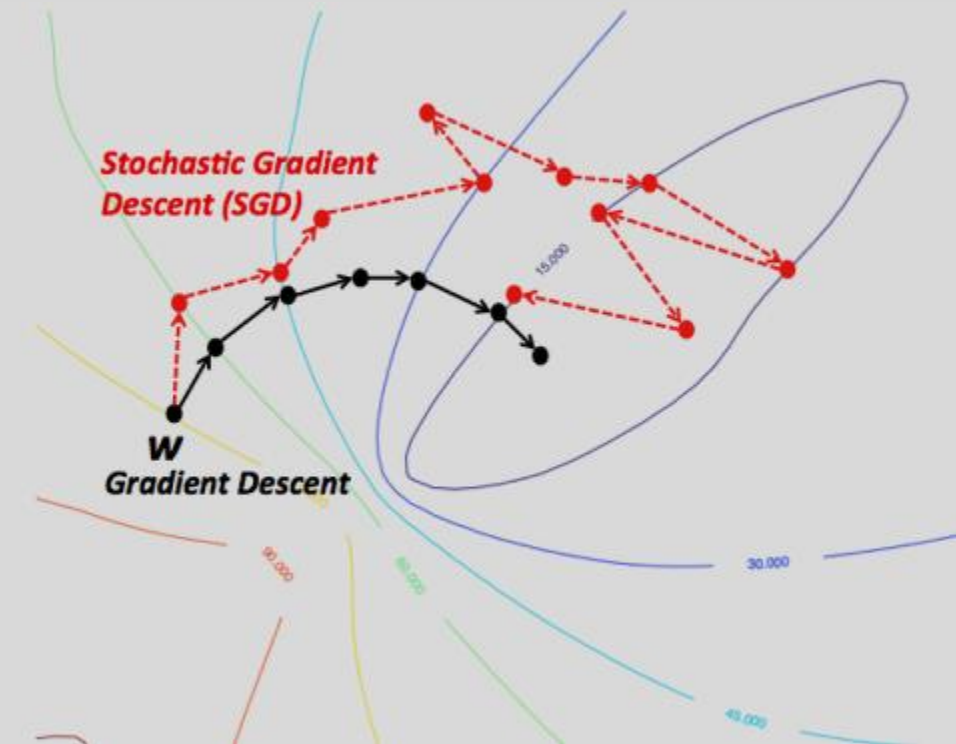
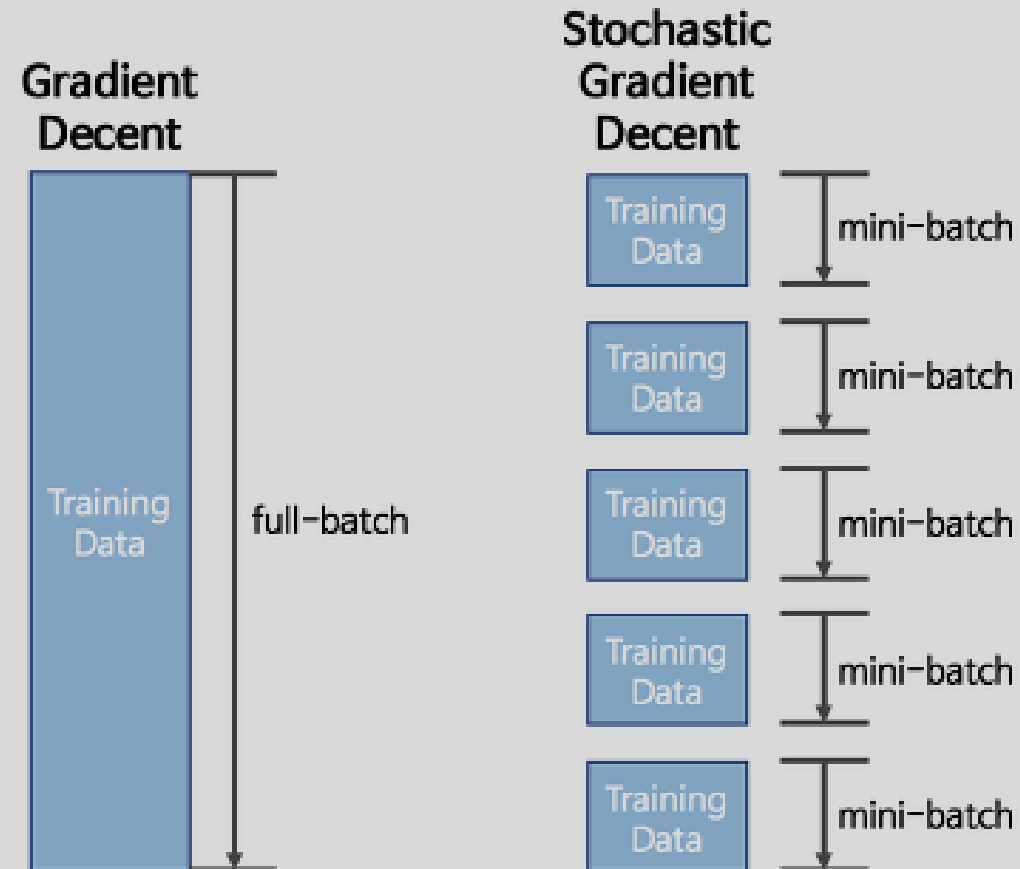
```
for X, y_true in train_loader:
```



# Stochastic Gradient Descent


- Gradient descent
  - Calculate for all data (takes large amount of time).
  - Go 1 optimal step.
- Stochastic gradient descent
  - Calculate gradients for partial data (takes small amount of time).
  - Go many non-globally-optimal steps, but converges.

# Stochastic Gradient Descent




# Implementing LeNet

```
def train(train_loader, model, criterion, optimizer):  
  
    model.train()  
    train_loss = 0  
    correct = 0  
  
    for X, y_true in train_loader:  
  
        optimizer.zero_grad()  
  
        y_hat = model(X)  
        loss = criterion(y_hat, y_true)  
  
        train_loss += loss.item()  
        pred = y_hat.argmax(dim=1, keepdim=True)  
        correct += pred.eq(y_true.view_as(pred)).sum().item()  
  
        loss.backward()  
        optimizer.step()  
  
    epoch_loss = train_loss / len(train_loader.dataset)  
    acc = correct / len(train_loader.dataset)  
  
    return model, optimizer, epoch_loss, acc
```



Training flag .train()



Same as before.

# Implementing LeNet

```
def test(test_loader, model, criterion):
```

```
    model.eval()  
    test_loss = 0  
    correct = 0
```


```
    for X, y_true in test_loader:
```

```
        y_hat = model(X)  
        loss = criterion(y_hat, y_true)
```


```
        test_loss += loss.item()  
        pred = y_hat.argmax(dim=1, keepdim=True)  
        correct += pred.eq(y_true.view_as(pred)).sum().item()
```

```
    epoch_loss = test_loss / len(test_loader.dataset)  
    acc = correct / len(test_loader.dataset)
```

```
    return model, epoch_loss, acc
```




Testing flag .eval()



Same as before, but without backward step.

# Implementing LeNet

```
def training_loop(model, criterion, optimizer, train_loader, test_loader, epochs, print_every=1):  
  
    train_losses = []  
    test_losses = []  
  
    for epoch in range(epochs):  
  
        model, optimizer, train_loss, train_acc = train(train_loader, model, criterion, optimizer)  
        train_losses.append(train_loss)  
  
        with torch.no_grad():  
            model, test_loss, test_acc = test(test_loader, model, criterion)  
            test_losses.append(test_loss)  
  
        if epoch % print_every == (print_every - 1):  
  
            print(f'Epoch: {epoch}\t'  
                  f'Train loss: {train_loss:.4f}\t'  
                  f'Test loss: {test_loss:.4f}\t'  
                  f'Train accuracy: {100 * train_acc:.2f}\t'  
                  f'Test accuracy: {100 * test_acc:.2f}')  
    return model, optimizer, (train_losses, test_losses)
```



No gradient calculation mode.

# Implementing LeNet

```
model = LeNet5(N_CLASSES)

optimizer = torch.optim.SGD(model.parameters(), lr=LEARNING_RATE)
criterion = nn.CrossEntropyLoss()

model, optimizer, _ = training_loop(model, criterion, optimizer, train_loader, test_loader, N_EPOCHS)
```

# Training using CPU

```
Epoch: 0 Train loss: 2.2817 Test loss: 2.2540 Train accuracy: 14.16 Test accuracy: 34.11
Epoch: 1 Train loss: 2.1948 Test loss: 2.0867 Train accuracy: 50.41 Test accuracy: 56.95
Epoch: 2 Train loss: 1.8551 Test loss: 1.5684 Train accuracy: 56.49 Test accuracy: 62.10
Epoch: 3 Train loss: 1.3052 Test loss: 1.0489 Train accuracy: 67.44 Test accuracy: 74.17
Epoch: 4 Train loss: 0.9017 Test loss: 0.7621 Train accuracy: 77.70 Test accuracy: 81.95
Epoch: 5 Train loss: 0.6930 Test loss: 0.6112 Train accuracy: 82.86 Test accuracy: 85.24
Epoch: 6 Train loss: 0.5783 Test loss: 0.5238 Train accuracy: 85.29 Test accuracy: 86.71
Epoch: 7 Train loss: 0.5090 Test loss: 0.4686 Train accuracy: 86.75 Test accuracy: 87.75
Epoch: 8 Train loss: 0.4633 Test loss: 0.4298 Train accuracy: 87.60 Test accuracy: 88.62
Epoch: 9 Train loss: 0.4306 Test loss: 0.4009 Train accuracy: 88.34 Test accuracy: 89.17
Epoch: 10 Train loss: 0.4056 Test loss: 0.3788 Train accuracy: 88.82 Test accuracy: 89.68
Epoch: 11 Train loss: 0.3855 Test loss: 0.3602 Train accuracy: 89.27 Test accuracy: 90.08
Epoch: 12 Train loss: 0.3686 Test loss: 0.3444 Train accuracy: 89.65 Test accuracy: 90.45
Epoch: 13 Train loss: 0.3540 Test loss: 0.3310 Train accuracy: 89.95 Test accuracy: 90.75
Epoch: 14 Train loss: 0.3407 Test loss: 0.3183 Train accuracy: 90.29 Test accuracy: 91.00
```



It takes about 15 minutes for training 15 epochs. – **slow!**

# Using GPUs in Colab

```
import torch
print(torch.cuda.is_available())
```

False

Colab does not support GPUs by default.



Select GPU.

```
import torch
print(torch.cuda.is_available())
```

True

Colab now supports GPUs!



# Implementing LeNet

```
def train(train_loader, model, criterion, optimizer, device):  
  
    model.train()  
    train_loss = 0  
    correct = 0  
  
    for X, y_true in train_loader:  
        optimizer.zero_grad()  
  
        X = X.to(device)  
        y_true = y_true.to(device)  
  
        y_hat = model(X)  
        loss = criterion(y_hat, y_true)  
  
        train_loss += loss.item()  
        pred = y_hat.argmax(dim=1, keepdim=True)  
        correct += pred.eq(y_true.view_as(pred)).sum().item()  
  
        loss.backward()  
        optimizer.step()  
  
    epoch_loss = train_loss / len(train_loader.dataset)  
    acc = correct / len(train_loader.dataset)  
  
    return model, optimizer, epoch_loss, acc
```

# Implementing LeNet

```
def test(test_loader, model, criterion, device):  
  
    model.eval()  
    test_loss = 0  
    correct = 0  
  
    for X, y_true in test_loader:  
  
        X = X.to(device)  
        y_true = y_true.to(device)  
  
        y_hat = model(X)  
        loss = criterion(y_hat, y_true)  
  
        test_loss += loss.item()  
        pred = y_hat.argmax(dim=1, keepdim=True)  
        correct += pred.eq(y_true.view_as(pred)).sum().item()  
  
    epoch_loss = test_loss / len(test_loader.dataset)  
    acc = correct / len(test_loader.dataset)  
  
    return model, epoch_loss, acc
```

# Implementing LeNet

```
def training_loop(model, criterion, optimizer, train_loader, test_loader, epochs, device, print_every=1):  
  
    best_loss = 1e10  
    train_losses = []  
    test_losses = []  
  
    for epoch in range(epochs):  
  
        model, optimizer, train_loss, train_acc = train(train_loader, model, criterion, optimizer, device)  
        train_losses.append(train_loss)  
  
        with torch.no_grad():  
            model, test_loss, test_acc = test(test_loader, model, criterion, device)  
            test_losses.append(test_loss)  
  
        if epoch % print_every == (print_every - 1):  
  
            print(f'Epoch: {epoch}\t'  
                  f'Train loss: {train_loss:.4f}\t'  
                  f'Test loss: {test_loss:.4f}\t'  
                  f'Train accuracy: {100 * train_acc:.2f}\t'  
                  f'Test accuracy: {100 * test_acc:.2f}')  
    return model, optimizer, (train_losses, test_losses)
```

# Implementing LeNet

```
DEVICE = 'cuda' if torch.cuda.is_available() else 'cpu'

model = LeNet5(N_CLASSES).to(DEVICE)

optimizer = torch.optim.SGD(model.parameters(), lr=LEARNING_RATE)
criterion = nn.CrossEntropyLoss()


model, optimizer, _ = training_loop(model, criterion, optimizer, train_loader, test_loader, N_EPOCHS,
DEVICE)
```


# Training using GPU

```
Epoch: 0 Train loss: 2.2817 Test loss: 2.2540 Train accuracy: 14.16 Test accuracy: 34.11
Epoch: 1 Train loss: 2.1948 Test loss: 2.0867 Train accuracy: 50.41 Test accuracy: 56.95
Epoch: 2 Train loss: 1.8551 Test loss: 1.5684 Train accuracy: 56.49 Test accuracy: 62.10
Epoch: 3 Train loss: 1.3052 Test loss: 1.0489 Train accuracy: 67.44 Test accuracy: 74.17
Epoch: 4 Train loss: 0.9017 Test loss: 0.7621 Train accuracy: 77.70 Test accuracy: 81.95
Epoch: 5 Train loss: 0.6930 Test loss: 0.6112 Train accuracy: 82.86 Test accuracy: 85.24
Epoch: 6 Train loss: 0.5783 Test loss: 0.5238 Train accuracy: 85.29 Test accuracy: 86.71
Epoch: 7 Train loss: 0.5090 Test loss: 0.4686 Train accuracy: 86.75 Test accuracy: 87.75
Epoch: 8 Train loss: 0.4633 Test loss: 0.4298 Train accuracy: 87.60 Test accuracy: 88.62
Epoch: 9 Train loss: 0.4306 Test loss: 0.4009 Train accuracy: 88.34 Test accuracy: 89.17
Epoch: 10 Train loss: 0.4056 Test loss: 0.3788 Train accuracy: 88.82 Test accuracy: 89.68
Epoch: 11 Train loss: 0.3855 Test loss: 0.3602 Train accuracy: 89.27 Test accuracy: 90.08
Epoch: 12 Train loss: 0.3686 Test loss: 0.3444 Train accuracy: 89.65 Test accuracy: 90.45
Epoch: 13 Train loss: 0.3540 Test loss: 0.3310 Train accuracy: 89.95 Test accuracy: 90.75
Epoch: 14 Train loss: 0.3407 Test loss: 0.3183 Train accuracy: 90.29 Test accuracy: 91.00
```

Same results but it only takes 1-2 mins. For 15 epochs. – **much faster!**

# Next class: AlexNet, VGG, GoogLeNet, ResNet





**Thank you!**

**UNIST**

**ULSAN NATIONAL INSTITUTE OF  
SCIENCE AND TECHNOLOGY**

**2 0 0 7**